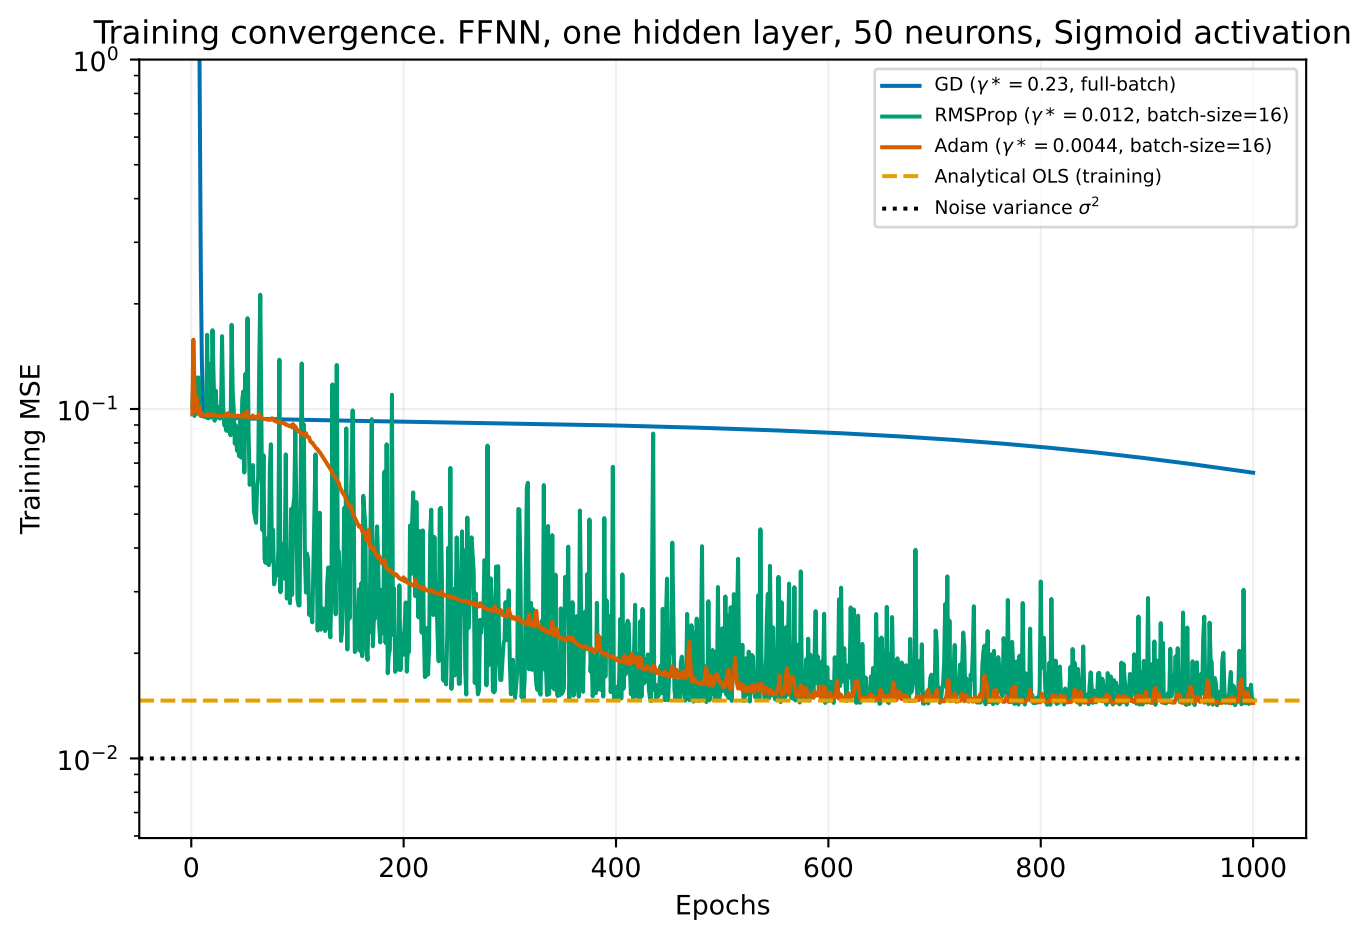
Training convergence. FFNN, one hidden layer, 50 neurons, Sigmoid activation
GD ($\gamma* = 0.23$, full-batch)
RMSProp ($\gamma* = 0.012$, batch-size=16)
Adam ($\gamma* = 0.0044$, batch-size=16)
Analytical OLS (training)
Noise variance σ^2
Training MSE
10^0
10^{-1}
10^{-2}
0
200
400
600
800
1000
Epochs